

DSP → Robot Policy Safety Translation

Objective: Show how caregiving reinforcement maps to robot constraints.

Mapping: Positive reinforcement = reward; Redirection = small penalties; Ethical constraint = hard penalty + termination.

Simulation: 7x7 gridworld; robot picks up object and reaches goal; human zone must be avoided.

Result: Baseline agent cuts through human zone; Constrained agent avoids it. Safety constraints change behavior at the cost of efficiency.

Tesla Implication: DSP-style reinforcement can become Optimus safety policies. Ethical constraints encoded as inviolable rules build trust.